

# Training a CNN

In this lab we are going to be working with a "[17 Category Flower Dataset](#)" from Visual Geometry Group of Oxford University. We will acquire the data, split it, train multiple models and do some visualizations.

```
In [1]: import numpy as np
import pandas as pd
import os
import shutil
```

## 1. Data aquisition

First, let's download the data from the webpage. You could have done it manually by going to the page, but we'll do it in the script.

```
In [2]: import urllib
```

```
In [3]: dataset_url = "https://www.robots.ox.ac.uk/~vgg/data/flowers/17/17flowers.tgz"
split_description_url = "https://www.robots.ox.ac.uk/~vgg/data/flowers/17/datasplits.mat"
#segmentation_ground_truth_url = "https://www.robots.ox.ac.uk/~vgg/data/flowers/17/trimaps.tgz"
readme_url = "https://www.robots.ox.ac.uk/~vgg/data/flowers/17/README.txt"
```

First, let's download the README file

```
In [4]: # create folder to store data
data_folder = "data/"
os.makedirs(data_folder, exist_ok=True)
```

```
In [5]: # Let's write a function to download data as we'll use multiple times
def get_file(file_url, target_folder=""):
    filename = os.path.basename(file_url)
    # express explicitly the filepath where data will be downloaded
    target_filepath = os.path.join(target_folder, filename)
    filepath, response = urllib.request.urlretrieve(file_url, target_filepath)

    return filepath, response
```

```
In [6]: # download readme file
readme_filepath, response = get_file(readme_url, data_folder)
```

```
In [7]: # Check out the README
with open(readme_filepath, 'r') as readme:
    text = readme.read()
    print(text)
```

17 Flower Category Database

-----

This set contains images of flowers belonging to 17 different categories. The images were acquired by searching the web and taking pictures. There are 80 images for each category.

The database was used in:

Nilsback, M-E. and Zisserman, A. A Visual Vocabulary for Flower Classification. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (2006) <http://www.robots.ox.ac.uk/~vgg/publications/papers/nilsback06.{pdf,ps.gz}>.

The datasplits used in this paper are specified in datasplits.mat

There are 3 separate splits. The results in the paper are averaged over the 3 splits. Each split has a training file (trn1, trn2, trn3), a validation file (val1, val2, val3) and a testfile (tst1, tst2 or tst3).

Segmentation Ground Truth

-----  
The ground truth is given for a subset of the images from 13 different categories.

More details can be found in:

Nilsback, M-E. and Zisserman, A. Delving into the whorl of flower segmentation. Proceedings of the British Machine Vision Conference (2007)  
[http://www.robots.ox.ac.uk/~vgg/publications/papers/nilsback06.\(pdf,ps.gz\).](http://www.robots.ox.ac.uk/~vgg/publications/papers/nilsback06.(pdf,ps.gz).)

The ground truth file also contains the file `imlist.mat`, which indicated which images in the original database that have been anotated.

Distance matrices  
-----

We provide two set of distance matrices:

1. `distancematrices17gcfeat06.mat`
  - Distance matrices using the same features and segmentation as detailed in:  
Nilsback, M-E. and Zisserman, A. A Visual Vocabulary for Flower Classification. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition(2006)  
<http://www.robots.ox.ac.uk/~vgg/publications/papers/nilsback06.{pdf,ps.gz}.>
2. `distancematrices17itfeat08.mat`
  - Distance matrices using the same features as described in:  
Nilsback, M-E. and Zisserman, A. Automated flower classification over a large number of classe S.  
Proceedings of the Indian Conference on Computer Vision, Graphics and Image Processing (2008)  
<http://www.robots.ox.ac.uk/~vgg/publications/papers/nilsback08.{pdf,ps.gz}.>  
and the iterative segmenation scheme detailed in  
Nilsback, M-E. and Zisserman, A. Delving into the whorl of flower segmentation. Proceedings of the British Machine Vision Conference (2007)  
[http://www.robots.ox.ac.uk/~vgg/publications/papers/nilsback06.\(pdf,ps.gz\).](http://www.robots.ox.ac.uk/~vgg/publications/papers/nilsback06.(pdf,ps.gz).)

Now, let's download the data

```
In [8]: # download the data
dataset_filepath, response = get_file(dataset_url, data_folder)
```

We have just downloaded a tar file. Let's unpack it.

```
In [9]: import tarfile
```

```
In [10]: with tarfile.open(dataset_filepath) as tar:
tar.extractall(path=data_folder)
```

What have we extracted?

```
In [11]: os.listdir(data_folder)
```

```
Out[11]: ['17flowers.tgz', 'datasplits.mat', 'jpg', 'README.txt', 'training_folder']
```

We see that a new folder named *jpg* has appeared.

```
In [12]: os.listdir(os.path.join(data_folder, 'jpg'))[:10]
```

```
Out[12]: ['files.txt',
'files.txt~',
'image_0001.jpg',
'image_0002.jpg',
'image_0003.jpg',
'image_0004.jpg',
'image_0005.jpg',
'image_0006.jpg',
'image_0007.jpg',
'image_0008.jpg']
```

This folder contains images of the dataset. But what about ground truth?

Based on the README, each class contains exactly 80 images. Quick check shows that images of one class are grouped together. We will use this fact later to group the images by class.

The split information was already provided with the dataset (otherwise we could have used `train_test_split` to obtain it)

```
In [13]: # download split file
split_filepath, response = get_file(split_description_url, data_folder)
```

```
In [14]: from scipy.io import loadmat
split = loadmat(split_filepath)
```

```
In [15]: split.keys()
```

```
Out[15]: dict_keys(['__header__', '__version__', '__globals__', 'trn1', 'trn2', 'trn3', 'tst1', 'tst2', 'tst3', 'val3', 'val2', 'val1'])
```

Let's use option 1 of train/val/test split:

```
In [16]: train = split["trn1"]
val = split["val1"]
test = split["tst1"]

print("""Train set contains {} files,
val set contains {} files,
and test set contains {} files""".format(train.shape[1], val.shape[1], test.shape[1]))
```

```
Train set contains 680 files,
val set contains 340 files,
and test set contains 340 files
```

## Excercise

Additional things to do:

- Check how many images we have downloaded.

```
In [17]: assert(17*80==len(train[0]) + len(val[0]) + len(test[0]))
print(f"We have {17*80} images")
```

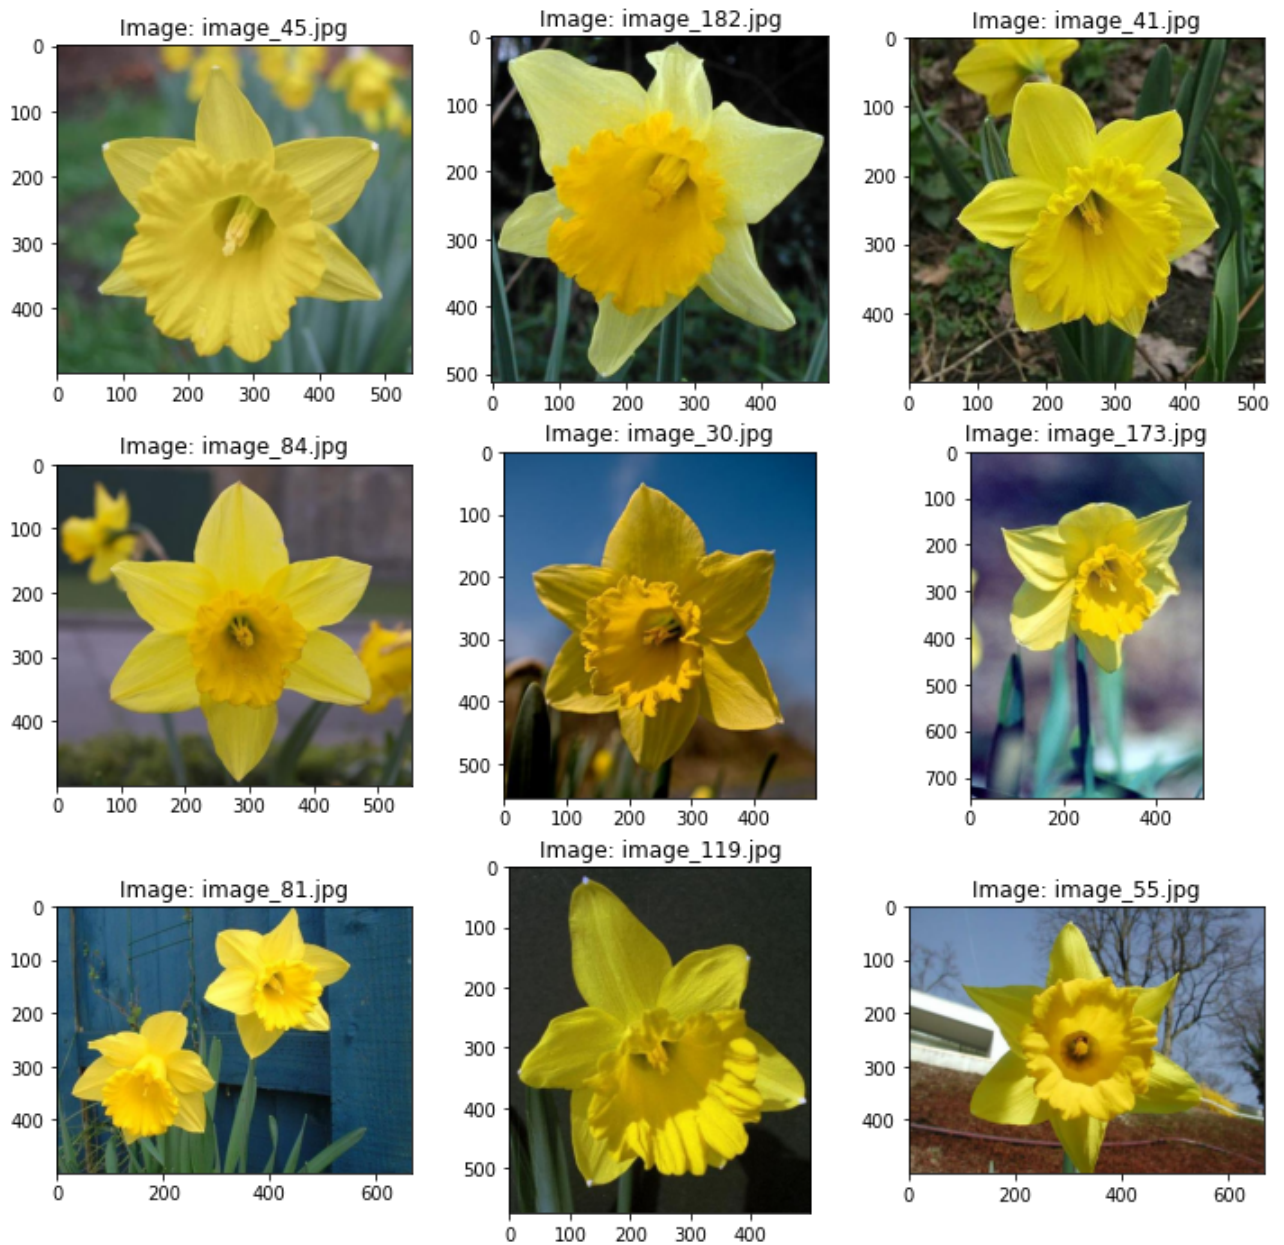
We have 1360 images

- Display some of the images.

```
In [18]: import matplotlib.image as mpimg
import matplotlib.pyplot as plt
img = mpimg.imread(os.path.join(data_folder, f'jpg/image_00{train[0][42]}.jpg'),)

path = os.path.join(data_folder, 'jpg/image_00')

indices = np.random.randint(0, 200, size=9)
fig, axes = plt.subplots(3, 3, figsize=(12, 12))
for ind, ax in enumerate(axes.flat):
    index = train[0][ind]
    if index < 10:
        imagepath = path + '0' + str(index) + '.jpg'
    else:
        imagepath = path + str(index) + '.jpg'
    img = mpimg.imread(imagepath)
    ax.imshow(img)
    ax.set_title(f"Image: image_{indices[ind]}.jpg")
```



- Are those color images?

-YES

- What are their shape?

```
In [19]: img.shape #example shape
```

```
Out[19]: (500, 666, 3)
```

- Are they all of the same shape?

```
In [20]: img2 = mpimg.imread(path+'47.jpg')
img2.shape
```

```
Out[20]: (760, 500, 3)
```

-NO, the shapes are different

## 2. Data regrouping

During the training with Keras for the simplicity we are going to be using `flow_from_dir` method of `ImageDataGenerator`. However, we'll need to organize data first in the specific manner: separate train, val, test sets, and put images of each class in a designated folder.

First, let's write a function to get a class name from file index. We'll use the fact that each class has 80 images, and they are grouped together by index.

```
In [21]: def get_image_class(file_index):
         image_class_idx = (int(file_index) - 1) // 80 + 1
         class_name = "{:02d}".format(image_class_idx)

         return class_name
```

Now let's rearrange the data

```
In [22]: from shutil import copy
```

```
In [23]: training_folder_name = "training_folder"
```

```
In [24]: for filename in os.listdir(os.path.join(data_folder, 'jpg')):
         if filename.endswith('jpg'):
             ### filename 'image_0936.jpg' --> file_index 936
             file_index = int(filename[6:10])
             true_class = get_image_class(file_index)
             if file_index in train:
                 split_folder = 'train'
             elif file_index in val:
                 split_folder = 'val'
             elif file_index in test:
                 split_folder = 'test'

             target_folder = os.path.join(data_folder, training_folder_name, split_folder, true_class)
             os.makedirs(target_folder, exist_ok=True)

             source_filepath = os.path.join(data_folder, 'jpg', filename)
             copy(source_filepath, target_folder)
         else:
             print(filename)
             print("Not a jpg file, skipping")
```

```
files.txt
Not a jpg file, skipping
files.txt~
Not a jpg file, skipping
```

## 3. CNN training

Now that we have prepared the data, we will be able to train a model.

### 3.1 Transfer learning

Let's do the [transfer learning](#) we have briefly discussed last time. We'll load one of the pretrained models from Keras library with [ImageNet weights](#).

#### Model preparation

```
In [25]: import tensorflow as tf
         import tensorflow.keras as keras
         from tensorflow.keras.preprocessing.image import ImageDataGenerator
         from tensorflow.keras import backend as K
```

```
import os
```

```
c:\users\eddie\kma_course\lib\site-packages\tensorflow\python\framework\dtypes.py:516: FutureWarnin
g: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, i
t will be understood as (type, (1,)) / '(1,)type'.
_np_qint8 = np.dtype(["qint8", np.int8, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorflow\python\framework\dtypes.py:517: FutureWarnin
g: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, i
t will be understood as (type, (1,)) / '(1,)type'.
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorflow\python\framework\dtypes.py:518: FutureWarnin
g: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, i
t will be understood as (type, (1,)) / '(1,)type'.
_np_qint16 = np.dtype(["qint16", np.int16, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorflow\python\framework\dtypes.py:519: FutureWarnin
g: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, i
t will be understood as (type, (1,)) / '(1,)type'.
_np_quint16 = np.dtype(["quint16", np.uint16, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorflow\python\framework\dtypes.py:520: FutureWarnin
g: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, i
t will be understood as (type, (1,)) / '(1,)type'.
_np_qint32 = np.dtype(["qint32", np.int32, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorflow\python\framework\dtypes.py:525: FutureWarnin
g: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, i
t will be understood as (type, (1,)) / '(1,)type'.
_np_resource = np.dtype(["resource", np.ubyte, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:541: Futur
eWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of n
umpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_qint8 = np.dtype(["qint8", np.int8, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:542: Futur
eWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of n
umpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:543: Futur
eWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of n
umpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_qint16 = np.dtype(["qint16", np.int16, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:544: Futur
eWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of n
umpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_quint16 = np.dtype(["quint16", np.uint16, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:545: Futur
eWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of n
umpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_qint32 = np.dtype(["qint32", np.int32, 1])
c:\users\eddie\kma_course\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:550: Futur
eWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of n
umpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_resource = np.dtype(["resource", np.ubyte, 1])
```

```
In [26]: # Reproducibility!
np.random.seed(42)
tf.set_random_seed(42)
```

```
In [27]: ## """"## GPU selection --> execute do only if you need to select a GPU / part of GPU

# os.environ["CUDA_VISIBLE_DEVICES"] = "0"

# #### Set session with share of GPU
# config_1 = tf.ConfigProto()
# # gpu_fraction_1 = float(os.environ.get('GPU_LIMIT_1', 0.45))
# # config_1.gpu_options.per_process_gpu_memory_fraction = gpu_fraction_1
# config_1.gpu_options.allow_growth = True

# sess_1 = tf.Session(config=config_1)
# sess_1.run(tf.global_variables_initializer())
# K.set_session(sess_1)
```

We'll be using VGG16 model. Together with weights, we'll also need a corresponding preprocessing function for the input images.

```
In [28]: from tensorflow.keras.applications.vgg16 import VGG16
        from tensorflow.keras.applications.vgg16 import preprocess_input as preprocess_input_vgg

        from tensorflow.keras.layers import Dense, Dropout, Flatten
        from tensorflow.keras.models import Model
```

```
In [29]: base_model = VGG16(include_top=False, weights='imagenet', input_shape = (224,224,3))
        base_model.summary()
```

WARNING:tensorflow:From c:\users\eddie\kma\_course\lib\site-packages\tensorflow\python\ops\init\_ops.py:1251: calling VarianceScaling.\_\_init\_\_ (from tensorflow.python.ops.init\_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor  
Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
<hr/>		
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
<hr/>		
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
<hr/>		
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
<hr/>		
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
<hr/>		
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
<hr/>		
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
<hr/>		
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
<hr/>		
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
<hr/>		
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
<hr/>		
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
<hr/>		
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

Note that we have downloaded only a convolution part of the neural network. Let's add some dense layers on top of it.

```
In [30]: ## freezing the layers
```



```
#for layer in base_model.layers:
#    layer.trainable = False
```

```
In [31]: nb_classes = 17
```

```
In [32]: flatten = Flatten()(base_model.output)
dropout_1 = Dropout(0.25)(flatten)
fc_1 = Dense(1000)(dropout_1)
dropout_2 = Dropout(0.5)(fc_1)
predictions = Dense(nb_classes, activation="softmax", name='predictions')(dropout_2)
```

```
In [33]: model = Model(base_model.input, predictions)
```

```
In [34]: model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
dropout (Dropout)	(None, 25088)	0
dense (Dense)	(None, 1000)	25089000
dropout_1 (Dropout)	(None, 1000)	0
predictions (Dense)	(None, 17)	17017
=====		
Total params: 39,820,705		
Trainable params: 39,820,705		



Non-trainable params: 0

---

## Model training parameters

```
In [35]: from tensorflow.keras import optimizers
```

```
In [36]: loss = 'categorical_crossentropy'
learning_rate = 0.001
optimizer = optimizers.SGD ## optimizers.SGD ## optimizers.RMSprop ## optimizers.Adagrad ## optimizers.Adam
metrics = ['accuracy']
```

```
In [37]: model.compile(loss=loss,
                      optimizer=optimizer(learning_rate),
                      metrics=metrics)
```

## Data preparation

```
In [38]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
In [39]: train_dir = os.path.join(data_folder, training_folder_name, "train")
val_dir = os.path.join(data_folder, training_folder_name, "val")
test_dir = os.path.join(data_folder, training_folder_name, "test")
```

```
In [40]: # we'll resize images in correspondance to network input size
image_size = (224,224)
```

```
In [41]: # apply some data augmentation
train_datagen = ImageDataGenerator(rotation_range=15,
                                   width_shift_range=0.2,
                                   height_shift_range=0.2,
                                   horizontal_flip=True,
                                   fill_mode='nearest',
                                   preprocessing_function=preprocess_input_vgg
                                   )

validation_datagen = ImageDataGenerator(preprocessing_function=preprocess_input_vgg) # for validation

train_batchsize = 30
val_batchsize = 30

# this function takes images from folders and feeds to Imagedatagenerator
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=image_size,
    batch_size=train_batchsize,
    class_mode='categorical')

validation_generator = validation_datagen.flow_from_directory(
    val_dir,
    target_size=image_size,
    batch_size=val_batchsize,
    class_mode='categorical',
    shuffle=False)
```

Found 680 images belonging to 17 classes.  
Found 340 images belonging to 17 classes.

## Model training

```
In [42]: epochs = 50
```

```
In [43]: nb_train_steps = train_generator.samples // train_generator.batch_size
nb_val_steps = validation_generator.samples // validation_generator.batch_size
```

```
In [44]: history = model.fit_generator(  
        train_generator,  
        steps_per_epoch=nb_train_steps,  
        epochs=epochs,  
        validation_data=validation_generator,  
        validation_steps=nb_val_steps,  
        verbose=1, #0  
    )
```

```
Epoch 1/50  
22/22 [=====] - 33s 2s/step - loss: 5.1333 - acc: 0.0523 - val_loss: 2.830  
7 - val_acc: 0.0758  
Epoch 2/50  
22/22 [=====] - 22s 1s/step - loss: 2.8348 - acc: 0.0723 - val_loss: 2.827  
7 - val_acc: 0.0727  
Epoch 3/50  
22/22 [=====] - 22s 1s/step - loss: 2.8360 - acc: 0.0600 - val_loss: 2.826  
5 - val_acc: 0.0424  
Epoch 4/50  
22/22 [=====] - 22s 1s/step - loss: 2.8332 - acc: 0.0462 - val_loss: 2.819  
6 - val_acc: 0.0788  
Epoch 5/50  
22/22 [=====] - 22s 1s/step - loss: 2.8287 - acc: 0.0585 - val_loss: 2.818  
7 - val_acc: 0.1000  
Epoch 6/50  
22/22 [=====] - 22s 1s/step - loss: 2.8298 - acc: 0.0600 - val_loss: 2.815  
7 - val_acc: 0.0818  
Epoch 7/50  
22/22 [=====] - 22s 1s/step - loss: 2.8261 - acc: 0.0892 - val_loss: 2.821  
4 - val_acc: 0.0697  
Epoch 8/50  
22/22 [=====] - 22s 1s/step - loss: 2.8184 - acc: 0.0954 - val_loss: 2.806  
9 - val_acc: 0.0879  
Epoch 9/50  
22/22 [=====] - 24s 1s/step - loss: 2.7998 - acc: 0.0800 - val_loss: 2.800  
9 - val_acc: 0.1061  
Epoch 10/50  
22/22 [=====] - 24s 1s/step - loss: 2.8023 - acc: 0.0938 - val_loss: 2.799  
0 - val_acc: 0.0818  
Epoch 11/50  
22/22 [=====] - 22s 1s/step - loss: 2.8003 - acc: 0.0924 - val_loss: 2.786  
8 - val_acc: 0.0848  
Epoch 12/50  
22/22 [=====] - 22s 1000ms/step - loss: 2.7828 - acc: 0.0859 - val_loss:  
2.7827 - val_acc: 0.1364  
Epoch 13/50  
22/22 [=====] - 23s 1s/step - loss: 2.7571 - acc: 0.1246 - val_loss: 2.748  
7 - val_acc: 0.1121  
Epoch 14/50  
22/22 [=====] - 23s 1s/step - loss: 2.7983 - acc: 0.0815 - val_loss: 2.760  
5 - val_acc: 0.1121  
Epoch 15/50  
22/22 [=====] - 22s 1s/step - loss: 2.7679 - acc: 0.1138 - val_loss: 2.750  
6 - val_acc: 0.1182  
Epoch 16/50  
22/22 [=====] - 22s 1s/step - loss: 2.7466 - acc: 0.1242 - val_loss: 2.737  
9 - val_acc: 0.1242  
Epoch 17/50  
22/22 [=====] - 22s 1s/step - loss: 2.7417 - acc: 0.1231 - val_loss: 2.735  
2 - val_acc: 0.1364  
Epoch 18/50  
22/22 [=====] - 22s 1s/step - loss: 2.7137 - acc: 0.1123 - val_loss: 2.721  
2 - val_acc: 0.1333  
Epoch 19/50  
22/22 [=====] - 22s 1s/step - loss: 2.7077 - acc: 0.1492 - val_loss: 2.740  
6 - val_acc: 0.1515  
Epoch 20/50  
22/22 [=====] - 22s 1s/step - loss: 2.7028 - acc: 0.1415 - val_loss: 2.654  
0 - val_acc: 0.1455  
Epoch 21/50  
22/22 [=====] - 22s 1s/step - loss: 2.6628 - acc: 0.1585 - val_loss: 2.628  
5 - val_acc: 0.1636
```

Epoch 22/50  
22/22 [=====] - 22s 1s/step - loss: 2.6113 - acc: 0.1831 - val\_loss: 2.565  
2 - val\_acc: 0.1818  
Epoch 23/50  
22/22 [=====] - 22s 1s/step - loss: 2.5779 - acc: 0.2077 - val\_loss: 2.540  
1 - val\_acc: 0.2000  
Epoch 24/50  
22/22 [=====] - 22s 1s/step - loss: 2.4792 - acc: 0.2231 - val\_loss: 2.462  
2 - val\_acc: 0.2455  
Epoch 25/50  
22/22 [=====] - 22s 1s/step - loss: 2.3669 - acc: 0.2600 - val\_loss: 2.496  
9 - val\_acc: 0.2606  
Epoch 26/50  
22/22 [=====] - 22s 1s/step - loss: 2.4049 - acc: 0.2415 - val\_loss: 2.267  
0 - val\_acc: 0.3030  
Epoch 27/50  
22/22 [=====] - 22s 1s/step - loss: 2.2561 - acc: 0.2815 - val\_loss: 2.022  
4 - val\_acc: 0.3636  
Epoch 28/50  
22/22 [=====] - 22s 1s/step - loss: 2.1497 - acc: 0.2923 - val\_loss: 2.008  
0 - val\_acc: 0.3758  
Epoch 29/50  
22/22 [=====] - 22s 1s/step - loss: 2.2245 - acc: 0.3108 - val\_loss: 1.823  
8 - val\_acc: 0.4182  
Epoch 30/50  
22/22 [=====] - 22s 1s/step - loss: 2.0949 - acc: 0.3646 - val\_loss: 1.964  
8 - val\_acc: 0.3636  
Epoch 31/50  
22/22 [=====] - 22s 1s/step - loss: 1.9734 - acc: 0.3646 - val\_loss: 1.907  
4 - val\_acc: 0.3788  
Epoch 32/50  
22/22 [=====] - 22s 1s/step - loss: 1.8449 - acc: 0.3985 - val\_loss: 1.584  
1 - val\_acc: 0.4788  
Epoch 33/50  
22/22 [=====] - 22s 1s/step - loss: 1.7389 - acc: 0.4136 - val\_loss: 1.918  
3 - val\_acc: 0.4030  
Epoch 34/50  
22/22 [=====] - 22s 1s/step - loss: 1.6190 - acc: 0.4738 - val\_loss: 1.561  
2 - val\_acc: 0.5000  
Epoch 35/50  
22/22 [=====] - 22s 1s/step - loss: 1.9837 - acc: 0.4062 - val\_loss: 1.499  
7 - val\_acc: 0.5152  
Epoch 36/50  
22/22 [=====] - 22s 997ms/step - loss: 1.4996 - acc: 0.5188 - val\_loss: 1.  
3761 - val\_acc: 0.5515  
Epoch 37/50  
22/22 [=====] - 22s 1s/step - loss: 1.4945 - acc: 0.5030 - val\_loss: 1.367  
6 - val\_acc: 0.5212  
Epoch 38/50  
22/22 [=====] - 22s 1s/step - loss: 1.5708 - acc: 0.4815 - val\_loss: 1.329  
5 - val\_acc: 0.5879  
Epoch 39/50  
22/22 [=====] - 22s 1s/step - loss: 1.4546 - acc: 0.5308 - val\_loss: 1.175  
6 - val\_acc: 0.6061  
Epoch 40/50  
22/22 [=====] - 22s 1s/step - loss: 1.2193 - acc: 0.6062 - val\_loss: 1.086  
8 - val\_acc: 0.6515  
Epoch 41/50  
22/22 [=====] - 22s 995ms/step - loss: 1.2045 - acc: 0.6125 - val\_loss: 1.  
0878 - val\_acc: 0.6697  
Epoch 42/50  
22/22 [=====] - 22s 1s/step - loss: 1.1462 - acc: 0.6212 - val\_loss: 1.018  
8 - val\_acc: 0.6697  
Epoch 43/50  
22/22 [=====] - 22s 1s/step - loss: 1.1313 - acc: 0.6462 - val\_loss: 0.990  
5 - val\_acc: 0.6697  
Epoch 44/50  
22/22 [=====] - 22s 1s/step - loss: 1.0796 - acc: 0.6385 - val\_loss: 1.051  
1 - val\_acc: 0.6364  
Epoch 45/50  
22/22 [=====] - 22s 1s/step - loss: 0.9715 - acc: 0.6723 - val\_loss: 0.888  
9 - val\_acc: 0.7121

```

Epoch 46/50
22/22 [=====] - 22s 1s/step - loss: 0.9197 - acc: 0.7138 - val_loss: 1.507
6 - val_acc: 0.5606
Epoch 47/50
22/22 [=====] - 22s 1s/step - loss: 0.9764 - acc: 0.6954 - val_loss: 0.880
8 - val_acc: 0.7121
Epoch 48/50
22/22 [=====] - 22s 1s/step - loss: 0.7899 - acc: 0.7379 - val_loss: 0.798
1 - val_acc: 0.7667
Epoch 49/50
22/22 [=====] - 22s 998ms/step - loss: 0.8021 - acc: 0.7312 - val_loss: 0.
8964 - val_acc: 0.7000
Epoch 50/50
22/22 [=====] - 22s 1s/step - loss: 0.7591 - acc: 0.7631 - val_loss: 0.804
1 - val_acc: 0.7818

```

```

In [45]: print('training acc.: ', history.history['acc'][-1])
         print('val acc.: ', (history.history['val_acc'])[-1])

```

```

training acc.: 0.7630769
val acc.: 0.7818182

```

```


In [46]: import matplotlib.pyplot as plt
         %matplotlib inline
         def plot_history(history):
             plt.figure()
             plt.xlabel('Epoch')
             plt.ylabel('Accuracy %')
             plt.plot(history.epoch, np.array(history.history['acc']),
                      label='Train Accuracy')
             plt.plot(history.epoch, np.array(history.history['val_acc']),
                      label='Val Accuracy')
             plt.legend()
             plt.ylim([0, 1])

```

```

In [47]: plot_history(history)

```



### Save model

```

In [48]: weights_folder = "weights"
         os.makedirs(weights_folder, exist_ok=True)

         model_name = 'vgg16_transfer_weights.h5'
         model_path = os.path.join(weights_folder, model_name)

```

```

In [49]: # uncomment to save model
         # model.save(model_path)

```

### Do the test on images

```

from tensorflow.keras.preprocessing import image

```

```
In [50]: from tensorflow.keras.models import load_model
```

```
In [51]: # model = load_model(model_path, compile=False)
```

#### Single image prediction

```
In [52]: #test_dir = "data/training_folder/test"  
#image_size = (224,224)
```

```
In [53]: class_idx = '08'  
image_name = os.listdir(os.path.join(test_dir, class_idx))[0]  
image_path = os.path.join(test_dir, class_idx, image_name)
```

```
In [54]: image_path
```

```
Out[54]: 'data/training_folder\\test\\08\\image_0564.jpg'
```

```
In [55]: # predicting image: getting the output vector  
img = image.load_img(image_path, target_size=image_size)  
img_array = image.img_to_array(img)  
img_expanded = np.expand_dims(img_array, axis=0)  
preprocessed_image = preprocess_input_vgg(img_expanded)  
  
pred = model.predict(preprocessed_image)  
print(pred)  
  
[[2.5455248e-01 1.2235704e-06 4.3765303e-07 2.1682472e-05 1.9703209e-06  
 9.0924817e-07 1.5998894e-05 7.3553097e-01 2.3544788e-05 3.9512652e-04  
 1.7409110e-06 3.0837407e-05 1.4433003e-04 2.6464569e-03 6.3373023e-03  
 4.6452965e-07 2.9449983e-04]]
```

```
In [56]: img_expanded.shape
```

```
Out[56]: (1, 224, 224, 3)
```

```
In [57]: img_array.shape
```

```
Out[57]: (224, 224, 3)
```

```
In [58]: classes = ["{:02d}".format(i) for i in range(1, 18)]  
pred_class_idx = np.argmax(pred, axis=1)  
classes[pred_class_idx[0]]
```

```
Out[58]: '08'
```

```
In [59]: pred[0][pred_class_idx]
```

```
Out[59]: array([0.735531], dtype=float32)
```

#### Multiple image predictions

```
In [60]: from sklearn.metrics import classification_report, confusion_matrix  
import seaborn as sns; sns.set()
```

```
In [61]: test_datagen = ImageDataGenerator(preprocessing_function=preprocess_input_vgg)
```

```
In [62]: test_generator = test_datagen.flow_from_directory(  
    test_dir,  
    target_size=image_size,  
    shuffle = False,  
    class_mode='categorical',  
    batch_size=1)  
  
filenames = test_generator.filenames
```

```
nb_samples = len(filenamees)

predict = model.predict_generator(test_generator, steps=nb_samples)
```

Found 340 images belonging to 17 classes.

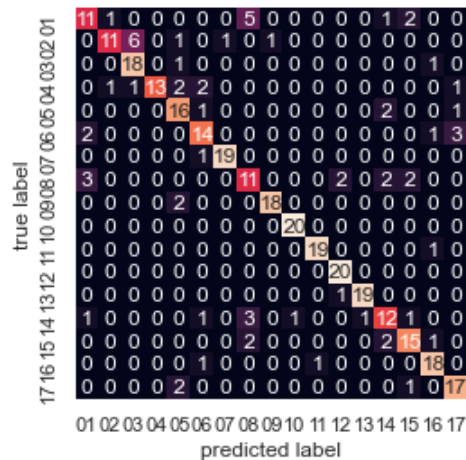
```
In [63]: predict.shape
```

```
Out[63]: (340, 17)
```

```
In [64]: y_pred = np.argmax(predict, axis=1)
print('Confusion Matrix')

mat = confusion_matrix(test_generator.classes, y_pred)
sns.heatmap(mat, square=True, annot=True, fmt='d', cbar=False,
            xticklabels=classes,
            yticklabels=classes)
plt.xlabel('predicted label')
plt.ylabel('true label');
```

Confusion Matrix



```
In [65]: #my results - Lesser batch size
#results heavily depend on seed - acc varies in range 0.2 - 0.9
print(classification_report(test_generator.classes + 1, y_pred + 1)) ## adding 1 to preserve the c
```

	precision	recall	f1-score	support
1	0.65	0.55	0.59	20
2	0.85	0.55	0.67	20
3	0.72	0.90	0.80	20
4	1.00	0.65	0.79	20
5	0.67	0.80	0.73	20
6	0.70	0.70	0.70	20
7	0.95	0.95	0.95	20
8	0.52	0.55	0.54	20
9	0.95	0.90	0.92	20
10	0.95	1.00	0.98	20
11	0.95	0.95	0.95	20
12	0.87	1.00	0.93	20
13	0.95	0.95	0.95	20
14	0.63	0.60	0.62	20
15	0.71	0.75	0.73	20
16	0.82	0.90	0.86	20
17	0.77	0.85	0.81	20
accuracy			0.80	340
macro avg	0.80	0.80	0.79	340
weighted avg	0.80	0.80	0.79	340

```
In [66]: ##original results
print(classification_report(test_generator.classes + 1, y_pred + 1)) ## adding 1 to preserve the c
```

	precision	recall	f1-score	support
1	0.65	0.55	0.59	20
2	0.85	0.55	0.67	20
3	0.72	0.90	0.80	20
4	1.00	0.65	0.79	20
5	0.67	0.80	0.73	20
6	0.70	0.70	0.70	20
7	0.95	0.95	0.95	20
8	0.52	0.55	0.54	20
9	0.95	0.90	0.92	20
10	0.95	1.00	0.98	20
11	0.95	0.95	0.95	20
12	0.87	1.00	0.93	20
13	0.95	0.95	0.95	20
14	0.63	0.60	0.62	20
15	0.71	0.75	0.73	20
16	0.82	0.90	0.86	20
17	0.77	0.85	0.81	20
accuracy				0.80
macro avg				0.80
weighted avg				0.80

Things to do

- Check some of incorrectly classified images
- Experiment with other models available in Keras
- Build your own network
- Optimize one or several training hyperparameters

## 3.2 Training enhancement

There are multiple ways to improve the quality of the model. Have a look at these papers that provide some heuristics for training a [classification](#) or [object detection](#) model.

I've done the following:

- used resnet50 with smaller images
- added custom LR Scheduler and EarlyStopping

```
In [67]: from tensorflow.keras.applications.resnet50 import ResNet50
         from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions
```

```
In [68]: base_model = ResNet50(include_top=False, weights='imagenet', input_shape = (112,112,3))
         # base_model.summary()

         #Note that we have downloaded only a convolution part of the neural network. Let's add some dense l

         nb_classes = 17

         flatten = Flatten()(base_model.output)
         dropout_1 = Dropout(0.25)(flatten)
         fc_1 = Dense(1000)(dropout_1)
         dropout_2 = Dropout(0.5)(fc_1)
         predictions = Dense(nb_classes, activation="softmax", name='predictions')(dropout_2)

         model = Model(base_model.input, predictions)

         # model.summary()
```

```
c:\users\eddie\kma_course\lib\site-packages\keras_applications\resnet50.py:265: UserWarning: The ou
tput shape of `ResNet50(include_top=False)` has been changed since Keras 2.2.0.
  warnings.warn('The output shape of `ResNet50(include_top=False)` '
```



```
In [69]: # we'll resize images in correspondance to network input size
image_size = (112,112)
```

```
In [70]: # apply some data augmentation
train_datagen = ImageDataGenerator(rotation_range=15,
                                   width_shift_range=0.2,
                                   height_shift_range=0.2,
                                   horizontal_flip=True,
                                   fill_mode='nearest',
                                   preprocessing_function=preprocess_input_vgg
                                   )

validation_datagen = ImageDataGenerator(preprocessing_function=preprocess_input_vgg) # for validation

train_batchsize = 30
val_batchsize = 30

# this function takes images from folders and feeds to Imagedatagenerator
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=image_size,
    batch_size=train_batchsize,
    class_mode='categorical')

validation_generator = validation_datagen.flow_from_directory(
    val_dir,
    target_size=image_size,
    batch_size=val_batchsize,
    class_mode='categorical',
    shuffle=False)

Found 680 images belonging to 17 classes.
Found 340 images belonging to 17 classes.
```

### Model training parameters

```
In [74]: loss = 'categorical_crossentropy'
learning_rate = 0.001
optimizer = optimizers.SGD ## optimizers.SGD ## optimizers.RMSprop ## optimizers.Adagrad ## optimizers.Adam
metrics = ['accuracy']
epochs = 50
model.compile(loss=loss,
              optimizer=optimizer(learning_rate),
              metrics=metrics)
```

```
In [72]: nb_train_steps = train_generator.samples // train_generator.batch_size
nb_val_steps = validation_generator.samples // validation_generator.batch_size
```

```
In [73]: def step_decay(epoch,lr):
    decay_rate = 0.9
    decay_step = 10
    if epoch% decay_step ==0 and epoch:
        return lr * decay_rate
    return lr
```

```
In [75]: history = model.fit_generator(
    train_generator,
    steps_per_epoch=nb_train_steps,
    epochs=epochs,
    validation_data=validation_generator,
    validation_steps=nb_val_steps,
    verbose=1, #0
    callbacks=[tf.keras.callbacks.EarlyStopping(patience=5,monitor='loss',restore_best_weights=True)]
)
```

Epoch 00001: LearningRateScheduler reducing learning rate to 0.0010000000474974513.  
Epoch 1/50

22/22 [=====] - 15s 691ms/step - loss: 5.6794 - acc: 0.2000 - val\_loss: 2.5075 - val\_acc: 0.4909

Epoch 00002: LearningRateScheduler reducing learning rate to 0.001000000474974513.  
Epoch 2/50  
22/22 [=====] - 7s 313ms/step - loss: 2.7803 - acc: 0.4600 - val\_loss: 2.3152 - val\_acc: 0.6273

Epoch 00003: LearningRateScheduler reducing learning rate to 0.001000000474974513.  
Epoch 3/50  
22/22 [=====] - 7s 310ms/step - loss: 1.8465 - acc: 0.6154 - val\_loss: 1.4151 - val\_acc: 0.7212

Epoch 00004: LearningRateScheduler reducing learning rate to 0.001000000474974513.  
Epoch 4/50  
22/22 [=====] - 7s 306ms/step - loss: 1.3209 - acc: 0.7123 - val\_loss: 1.1988 - val\_acc: 0.7576

Epoch 00005: LearningRateScheduler reducing learning rate to 0.001000000474974513.  
Epoch 5/50  
22/22 [=====] - 7s 300ms/step - loss: 0.9997 - acc: 0.7615 - val\_loss: 0.8386 - val\_acc: 0.8152

Epoch 00006: LearningRateScheduler reducing learning rate to 0.001000000474974513.  
Epoch 6/50  
22/22 [=====] - 7s 303ms/step - loss: 0.8266 - acc: 0.7909 - val\_loss: 0.7991 - val\_acc: 0.8303

Epoch 00007: LearningRateScheduler reducing learning rate to 0.001000000474974513.  
Epoch 7/50  
22/22 [=====] - 7s 298ms/step - loss: 0.7493 - acc: 0.8062 - val\_loss: 0.7146 - val\_acc: 0.8303

Epoch 00008: LearningRateScheduler reducing learning rate to 0.001000000474974513.  
Epoch 8/50  
22/22 [=====] - 7s 301ms/step - loss: 0.6125 - acc: 0.8431 - val\_loss: 0.6688 - val\_acc: 0.8515

Epoch 00009: LearningRateScheduler reducing learning rate to 0.001000000474974513.  
Epoch 9/50  
22/22 [=====] - 7s 301ms/step - loss: 0.6446 - acc: 0.8364 - val\_loss: 0.6411 - val\_acc: 0.8606

Epoch 00010: LearningRateScheduler reducing learning rate to 0.001000000474974513.  
Epoch 10/50  
22/22 [=====] - 7s 299ms/step - loss: 0.3959 - acc: 0.8859 - val\_loss: 0.5546 - val\_acc: 0.8636

Epoch 00011: LearningRateScheduler reducing learning rate to 0.000900000427477062.  
Epoch 11/50  
22/22 [=====] - 7s 310ms/step - loss: 0.3257 - acc: 0.9000 - val\_loss: 0.6239 - val\_acc: 0.8303

Epoch 00012: LearningRateScheduler reducing learning rate to 0.000900000427477062.  
Epoch 12/50  
22/22 [=====] - 7s 299ms/step - loss: 0.4195 - acc: 0.8831 - val\_loss: 0.5133 - val\_acc: 0.8636

Epoch 00013: LearningRateScheduler reducing learning rate to 0.000900000427477062.  
Epoch 13/50  
22/22 [=====] - 7s 299ms/step - loss: 0.3431 - acc: 0.8969 - val\_loss: 0.5840 - val\_acc: 0.8606

Epoch 00014: LearningRateScheduler reducing learning rate to 0.000900000427477062.  
Epoch 14/50  
22/22 [=====] - 7s 305ms/step - loss: 0.2808 - acc: 0.9200 - val\_loss: 0.5779 - val\_acc: 0.8727

Epoch 00015: LearningRateScheduler reducing learning rate to 0.000900000427477062.  
Epoch 15/50  
22/22 [=====] - 7s 302ms/step - loss: 0.3569 - acc: 0.8938 - val\_loss: 0.5916 - val\_acc: 0.8636

Epoch 00016: LearningRateScheduler reducing learning rate to 0.0009000000427477062.  
Epoch 16/50  
22/22 [=====] - 7s 306ms/step - loss: 0.2395 - acc: 0.9277 - val\_loss: 0.5232 - val\_acc: 0.8727

Epoch 00017: LearningRateScheduler reducing learning rate to 0.0009000000427477062.  
Epoch 17/50  
22/22 [=====] - 7s 306ms/step - loss: 0.2294 - acc: 0.9400 - val\_loss: 0.5014 - val\_acc: 0.8758

Epoch 00018: LearningRateScheduler reducing learning rate to 0.0009000000427477062.  
Epoch 18/50  
22/22 [=====] - 7s 309ms/step - loss: 0.1513 - acc: 0.9515 - val\_loss: 0.5535 - val\_acc: 0.8818

Epoch 00019: LearningRateScheduler reducing learning rate to 0.0009000000427477062.  
Epoch 19/50  
22/22 [=====] - 7s 304ms/step - loss: 0.2212 - acc: 0.9462 - val\_loss: 0.5148 - val\_acc: 0.8788

Epoch 00020: LearningRateScheduler reducing learning rate to 0.0009000000427477062.  
Epoch 20/50  
22/22 [=====] - 7s 304ms/step - loss: 0.1786 - acc: 0.9492 - val\_loss: 0.5076 - val\_acc: 0.8879

Epoch 00021: LearningRateScheduler reducing learning rate to 0.0008100000384729356.  
Epoch 21/50  
22/22 [=====] - 7s 304ms/step - loss: 0.1766 - acc: 0.9569 - val\_loss: 0.5503 - val\_acc: 0.8848

Epoch 00022: LearningRateScheduler reducing learning rate to 0.0008100000559352338.  
Epoch 22/50  
22/22 [=====] - 7s 313ms/step - loss: 0.1297 - acc: 0.9585 - val\_loss: 0.5742 - val\_acc: 0.8818

Epoch 00023: LearningRateScheduler reducing learning rate to 0.0008100000559352338.  
Epoch 23/50  
22/22 [=====] - 7s 303ms/step - loss: 0.1512 - acc: 0.9554 - val\_loss: 0.4994 - val\_acc: 0.8939

Epoch 00024: LearningRateScheduler reducing learning rate to 0.0008100000559352338.  
Epoch 24/50  
22/22 [=====] - 7s 298ms/step - loss: 0.2205 - acc: 0.9400 - val\_loss: 0.5376 - val\_acc: 0.8909

Epoch 00025: LearningRateScheduler reducing learning rate to 0.0008100000559352338.  
Epoch 25/50  
22/22 [=====] - 7s 297ms/step - loss: 0.1690 - acc: 0.9554 - val\_loss: 0.5475 - val\_acc: 0.8939

Epoch 00026: LearningRateScheduler reducing learning rate to 0.0008100000559352338.  
Epoch 26/50  
22/22 [=====] - 7s 300ms/step - loss: 0.1406 - acc: 0.9530 - val\_loss: 0.5203 - val\_acc: 0.9030

Epoch 00027: LearningRateScheduler reducing learning rate to 0.0008100000559352338.  
Epoch 27/50  
22/22 [=====] - 7s 298ms/step - loss: 0.1282 - acc: 0.9625 - val\_loss: 0.5081 - val\_acc: 0.9121

Epoch 00028: LearningRateScheduler reducing learning rate to 0.0008100000559352338.  
Epoch 28/50  
22/22 [=====] - 7s 298ms/step - loss: 0.1756 - acc: 0.9538 - val\_loss: 0.5260 - val\_acc: 0.8939

Epoch 00029: LearningRateScheduler reducing learning rate to 0.0008100000559352338.  
Epoch 29/50  
22/22 [=====] - 7s 298ms/step - loss: 0.1757 - acc: 0.9538 - val\_loss: 0.5174 - val\_acc: 0.8939

Epoch 00030: LearningRateScheduler reducing learning rate to 0.0008100000559352338.

```

Epoch 30/50
22/22 [=====] - 7s 298ms/step - loss: 0.1500 - acc: 0.9569 - val_loss: 0.5045 - val_acc: 0.8939

Epoch 00031: LearningRateScheduler reducing learning rate to 0.0007290000503417104.
Epoch 31/50
22/22 [=====] - 7s 304ms/step - loss: 0.0891 - acc: 0.9712 - val_loss: 0.4930 - val_acc: 0.9000

Epoch 00032: LearningRateScheduler reducing learning rate to 0.0007290000794455409.
Epoch 32/50
22/22 [=====] - 7s 301ms/step - loss: 0.1509 - acc: 0.9600 - val_loss: 0.5186 - val_acc: 0.9000

Epoch 00033: LearningRateScheduler reducing learning rate to 0.0007290000794455409.
Epoch 33/50
22/22 [=====] - 7s 302ms/step - loss: 0.1219 - acc: 0.9688 - val_loss: 0.4956 - val_acc: 0.9030

Epoch 00034: LearningRateScheduler reducing learning rate to 0.0007290000794455409.
Epoch 34/50
22/22 [=====] - 7s 310ms/step - loss: 0.0617 - acc: 0.9803 - val_loss: 0.4989 - val_acc: 0.9030

Epoch 00035: LearningRateScheduler reducing learning rate to 0.0007290000794455409.
Epoch 35/50
22/22 [=====] - 7s 309ms/step - loss: 0.1295 - acc: 0.9615 - val_loss: 0.5475 - val_acc: 0.8939

Epoch 00036: LearningRateScheduler reducing learning rate to 0.0007290000794455409.
Epoch 36/50
22/22 [=====] - 7s 312ms/step - loss: 0.1150 - acc: 0.9662 - val_loss: 0.5194 - val_acc: 0.9000

Epoch 00037: LearningRateScheduler reducing learning rate to 0.0007290000794455409.
Epoch 37/50
22/22 [=====] - 7s 304ms/step - loss: 0.1354 - acc: 0.9631 - val_loss: 0.5153 - val_acc: 0.9000

Epoch 00038: LearningRateScheduler reducing learning rate to 0.0007290000794455409.
Epoch 38/50
22/22 [=====] - 7s 304ms/step - loss: 0.0663 - acc: 0.9800 - val_loss: 0.5020 - val_acc: 0.8939

Epoch 00039: LearningRateScheduler reducing learning rate to 0.0007290000794455409.
Epoch 39/50
22/22 [=====] - 8s 341ms/step - loss: 0.0911 - acc: 0.9692 - val_loss: 0.4815 - val_acc: 0.8970

```

```
In [76]: print('training acc.: ', history.history['acc'][-1])
        print('val acc.: ', (history.history['val_acc'])[-1])
```


```

training acc.: 0.9692308
val acc.: 0.8969697

```

```
In [77]: import matplotlib.pyplot as plt
        %matplotlib inline
        def plot_history(history):
            plt.figure()
            plt.xlabel('Epoch')
            plt.ylabel('Accuracy %')
            plt.plot(history.epoch, np.array(history.history['acc']),
                     label='Train Accuracy')
            plt.plot(history.epoch, np.array(history.history['val_acc']),
                     label='Val Accuracy')
            plt.legend()
            plt.ylim([0, 1])
```

```
In [78]: plot_history(history)
```



### 3.2.1 Data augmentation

Better data augmentation can easily give a boost to a model. Some of the useful tools include

- [imgaug](#)
- [mixup](#)

In [76]: `!pip install imgaug`

```
Collecting imgaug
  Downloading imgaug-0.4.0-py2.py3-none-any.whl (948 kB)
Collecting opencv-python
  Downloading opencv_python-4.4.0.46-cp37-cp37m-win_amd64.whl (33.5 MB)
Requirement already satisfied: numpy>=1.15 in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from imgaug) (1.19.4)
Requirement already satisfied: matplotlib in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from imgaug) (3.3.3)
Collecting scikit-image>=0.14.2
  Downloading scikit_image-0.17.2-cp37-cp37m-win_amd64.whl (11.5 MB)
Requirement already satisfied: Pillow in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from imgaug) (8.0.1)
Collecting imageio
  Downloading imageio-2.9.0-py3-none-any.whl (3.3 MB)
Requirement already satisfied: six in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from imageio) (1.15.0)
Requirement already satisfied: scipy in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from imageio) (1.5.4)
Collecting Shapely
  Downloading Shapely-1.7.1-cp37-cp37m-win_amd64.whl (1.0 MB)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.3 in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from Shapely) (2.4.7)
Requirement already satisfied: python-dateutil>=2.1 in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from Shapely) (2.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from Shapely) (1.3.1)
Requirement already satisfied: cycler>=0.10 in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from Shapely) (0.10.0)
Collecting tifffile>=2019.7.26
  Downloading tifffile-2020.10.1-py3-none-any.whl (152 kB)
Collecting PyWavelets>=1.1.1
  Downloading PyWavelets-1.1.1-cp37-cp37m-win_amd64.whl (4.2 MB)
Collecting networkx>=2.0
  Downloading networkx-2.5-py3-none-any.whl (1.6 MB)
Requirement already satisfied: decorator>=4.3.0 in c:\users\eddie\appdata\local\programs\python\python37\lib\site-packages (from networkx>=2.0->scikit-image>=0.14.2->imgaug) (4.4.2)
Installing collected packages: opencv-python, tifffile, imageio, PyWavelets, networkx, scikit-image, Shapely, imgaug
Successfully installed PyWavelets-1.1.1 Shapely-1.7.1 imageio-2.9.0 imgaug-0.4.0 networkx-2.5 opencv-python-4.4.0.46 scikit-image-0.17.2 tifffile-2020.10.1
```

In [ ]: `from imgaug import augmenters as iaa`

```
seq = iaa.Sequential([
    iaa.Crop(px=(0,32)),
    iaa.FlipLR(0.5),
    iaa.GaussianBlur(sigma=(0,4.4))
])
base_model = ResNet50(include_top=False, weights='imagenet', input_shape = (112,112,3))
model = Model(base_model.input, predictions)
```

### 3.2.2 Learning rate scheduling and early stopping criteria

In Keras learning rate scheduling and early stopping criteria can be implemented using [Callbacks](#). In particular, the following are quite useful: LearningRateScheduler, ReduceLROnPlateau, EarlyStopping, CSVLogger, ModelCheckpoint.

In [ ]:

In [ ]:

In [ ]: